

Use the kubectl create command to create a Deployment that manages a Pod. The Pod runs a Container based on the provided Docker image.

View the Deployment

```
PS C:\Users\LENOVO> kubectl create deployment hello-node --image=registry.k8s.io/e2e-test-images/agnhost:2.53 -- /agnhost netexec --http-port=8080
deployment.apps/hello-node created
PS C:\Users\LENOVO> kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hello-node	1/1	1	1	15s

View cluster events

```
PS C:\Users\LENOVO> kubectl get events
```

LAST SEEN	TYPE	REASON	OBJECT
	MESSAGE		
5m13s	Normal	Scheduled	pod/hello-node-6c9b5f4b59
-pkb9v	Successfully assigned default/hello-node-6c9b5f4b59-pkb9v to minikube		
5m13s	Normal	Pulling	pod/hello-node-6c9b5f4b59
-pkb9v	Pulling image "registry.k8s.io/e2e-test-images/agnhost:2.53"		
5m2s	Normal	Pulled	pod/hello-node-6c9b5f4b59
-pkb9v	Successfully pulled image "registry.k8s.io/e2e-test-images/agnhost:2.53" in 10.903s (10.903s including waiting). Image size: 139374622 bytes.		
5m2s	Normal	Created	pod/hello-node-6c9b5f4b59
-pkb9v	Created container: agnhost		
5m1s	Normal	Started	pod/hello-node-6c9b5f4b59
-pkb9v	Started container agnhost		
5m13s	Normal	SuccessfulCreate	replicaset/hello-node-6c9b5f4b59
b5f4b59	Created pod: hello-node-6c9b5f4b59-pkb9v		
5m13s	Normal	ScalingReplicaSet	deployment/hello-node
	Scaled up replica set hello-node-6c9b5f4b59 from 0 to 1		
14m	Normal	Starting	node/minikube
	Starting kubelet.		
14m	Normal	NodeHasSufficientMemory	node/minikube
	Node minikube status is now: NodeHasSufficientMemory		
14m	Normal	NodeHasNoDiskPressure	node/minikube
	Node minikube status is now: NodeHasNoDiskPressure		
14m	Normal	Starting	node/minikube
	Starting kubelet.		
14m	Normal	NodeAllocatableEnforced	node/minikube
	Updated Node Allocatable limit across pods		

View the kubectl configuration:

```

PS C:\Users\LENOVO> kubectl config view
apiVersion: v1
clusters:
  provider: minikube.sigs.k8s.io
  version: v1.37.0
  name: context_info
  namespace: default
  user: minikube
  name: minikube
current-context: minikube
kind: Config
preferences: {}
users:
- name: minikube
  user:
    client-certificate: C:\Users\LENOVO\.minikube\profiles\minikube\client.crt
    client-key: C:\Users\LENOVO\.minikube\profiles\minikube\client.key

```

```

PS C:\Users\LENOVO> kubectl config view
apiVersion: v1
  version: v1.37.0
  name: cluster_info
  server: https://127.0.0.1:53473
  name: minikube
contexts:
- context:
  cluster: minikube
  extensions:
  - extension:
    last-update: Mon, 29 Sep 2025 14:58:01 CST
    provider: minikube.sigs.k8s.io
    version: v1.37.0
    name: context_info
    namespace: default
    user: minikube
  name: minikube
current-context: minikube
kind: Config
preferences: {}
users:
- name: minikube
  user:
    client-certificate: C:\Users\LENOVO\.minikube\profiles\minikube\client.crt
    client-key: C:\Users\LENOVO\.minikube\profiles\minikube\client.key

```

View application logs for a container in a pod (replace pod name with the one you got from `kubectl get pods`).

Create a Service

Expose the Pod to the public internet using the kubectl expose command:

View the Service you created:

```
PS C:\Users\LENOVO> kubectl expose deployment hello-node --type=LoadBalancer --port=8080
service/hello-node exposed
PS C:\Users\LENOVO> kubectl get services
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)
hello-node    LoadBalancer 10.97.9.167    <pending>      8080:32706/TCP
kubernetes    ClusterIP     10.96.0.1     <none>         443/TCP
PS C:\Users\LENOVO> minikube service hello-node
```


NAMESPACE	NAME	TARGET PORT	URL
default	hello-node	8080	http://192.168.49.2:32706

* 为服务 hello-node 启动隧道。 /

NAMESPACE	NAME	TARGET PORT	URL
default	hello-node		http://127.0.0.1:57205

* 为服务 hello-node 启动隧道。
* 正通过默认浏览器打开服务 default/hello-node...
! 因为你正在使用 windows 上的 Docker 驱动程序, 所以需要打开终端才能运行它。



Run the command to opens up a browser windoe



Enable addons

List the currently supported addons:

```
PS C:\Users\LENOVO> minikube addons list
```

ADDON NAME	PROFILE	STATUS	MAINTAINER
ambassador	minikube	disabled	3rd party (Ambassador)
amd-gpu-device-plugin	minikube	disabled	3rd party (AMD)
auto-pause	minikube	disabled	minikube
cloud-spanner	minikube	disabled	Google
csi-hostpath-driver	minikube	disabled	Kubernetes
dashboard	minikube	enabled 	Kubernetes
default-storageclass	minikube	enabled 	Kubernetes
efk	minikube	disabled	3rd party (Elastic)
freshpod	minikube	disabled	Google
gcp-auth	minikube	disabled	Google
gvisor	minikube	disabled	minikube
headlamp	minikube	disabled	3rd party (kinvolk.io)
inaccel	minikube	disabled	3rd party (InAccelerator [info@inaccel.com])
ingress	minikube	disabled	Kubernetes

View the Pod and Service you created by installing that addon

```
PS C:\Users\LENOVO> kubectl get pod,svc -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
pod/coredns-66bc5c9577-88x4z	1/1	Running	0	28m
pod/etcd-minikube	1/1	Running	0	29m
pod/kube-apiserver-minikube	1/1	Running	0	29m
pod/kube-controller-manager-minikube	1/1	Running	0	29m
pod/kube-proxy-2ltjl	1/1	Running	0	28m
pod/kube-scheduler-minikube	1/1	Running	0	29m
pod/metrics-server-85b7d694d7-5gkxb	0/1	Running	0	16s
pod/storage-provisioner	1/1	Running	0	29m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT
service/kube-dns	ClusterIP	10.96.0.10	<none>	53/U
service/metrics-server	ClusterIP	10.108.239.188	<none>	443/TCP

Check the output from metrics-server

```
PS C:\Users\LENOVO> kubectl top pods
error: Metrics API not available
PS C:\Users\LENOVO> kubectl top pods
```

NAME	CPU(cores)	MEMORY(bytes)
hello-node-6c9b5f4b59-pkb9v	1m	7Mi

Disable metrics-server

Now you can clean up the resources you created in your cluster:

Stop the Minikube cluster

```
PS C:\Users\LENOVO> minikube addons disable metrics-server
* 'metrics-server' 插件已被禁用
PS C:\Users\LENOVO> kubectl delete service hello-node
service "hello-node" deleted
PS C:\Users\LENOVO> kubectl delete deployment hello-node
deployment.apps "hello-node" deleted
PS C:\Users\LENOVO> minikube stop
* 正在停止节点 "minikube" ...
* 正在通过 SSH 关闭"minikube"...
* 1 个节点已停止。
```

Deploy an app

Let's deploy our first app on Kubernetes with the `kubectl create deployment` command. We need to provide the deployment name and app image location (include the full repository url for images hosted outside Docker Hub).

```
> kubectl create deployment kubernetes-bootcamp
--image=gcr.io/google-samples/kubernetes-bootcamp:v1

✓ run_shell_command kubectl create deployment kubernetes-bootcamp --image=gcr.io/googl...
deployment.apps/kubernetes-bootcamp created
```

list the deployments

```
PS C:\Users\LENOVO> kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 1/1     1            1           2m29s
```

View the app

The `kubectl proxy` command can create a proxy that will forward communications into the cluster-wide, private network. The proxy can be terminated by pressing control-C and won't show any output while it's running.

```
PS C:\Users\LENOVO> kubectl proxy
Starting to serve on 127.0.0.1:8001
```

You can see all those APIs hosted through the proxy endpoint. For example, we can query the version directly through the API using the `curl` command:

```
PS C:\Users\LENOVO> curl http://localhost:8001/version

StatusCode      : 200
StatusDescription : OK
Content         : {
  "major": "1",
  "minor": "34",
  Audit-Id: e8b353a7-21fd-48d4-a6a3-72aba62dd059
  X-Kubernetes-Pf-Flowschema-Uid: 92d88f97-842b-42b9-964c-d75ba27f2171
  X-Kubernetes-Pf-Prioritylevel-Uid: 81772df0-d503-4528-b0e9-b4d111...
  Forms      : {}
  Headers    : {[Audit-Id, e8b353a7-21fd-48d4-a6a3-72aba62dd059], [X-Kubernetes-Pf-Flowschema-Uid,
92d88f97-842b-42b9-964c-d75ba27f2171], [X-Kubernetes-Pf-Prioritylevel-Uid, 81772df0-
d503-4528-b0e9-b4d111f835f8], [Content-Length, 379]...}
  Images     : {}
  InputFields : {}
  Links      : {}
  ParsedHtml  : mshtml.HTMLDocumentClass
  RawContentLength : 379
}
```

The API server will automatically create an endpoint for each pod, based on the pod name,

that is also accessible through the proxy.

First get the Pod name, and store it in the environment variable `POD_NAME`.

```
PS C:\Users\LENOVO> echo Name of the Pod: $POD_NAME
Name
of
the
Pod:
```

access the Pod through the proxied API



Services and Labels

use the `kubectl get` command and look for existing Pods:

```
PS C:\Users\LENOVO> kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-658f6cbd58-f967h 1/1     Running   0           59m
```

list the current Services from our cluster:

```
PS C:\Users\LENOVO> kubectl get services
NAME         TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
kubernetes   ClusterIP   10.96.0.1    <none>        443/TCP    11d
```

To expose the deployment to external traffic, we'll use the `kubectl expose` command with the `--type=NodePort` option:

```
PS C:\Users\LENOVO> kubectl expose deployment/kubernetes-bootcamp --type="NodePort" --port 8080
service/kubernetes-bootcamp exposed
```

To find out what port was opened externally (for the type: NodePort Service) we'll run the describe service subcommand:

```
PS C:\Users\LENOVO> kubectl describe services/kubernetes-bootcamp
Name:                kubernetes-bootcamp
Namespace:           default
Labels:              app=kubernetes-bootcamp
Annotations:         <none>
Selector:            app=kubernetes-bootcamp
Type:               NodePort
IP Family Policy:    SingleStack
IP Families:         IPv4
IP:                 10.107.83.183
IPs:                10.107.83.183
Port:               <unset> 8080/TCP
TargetPort:         8080/TCP
NodePort:           <unset> 30960/TCP
Endpoints:          10.244.0.22:8080
Session Affinity:    None
External Traffic Policy: Cluster
Internal Traffic Policy: Cluster
Events:             <none>
```

Create an environment variable called NODE_PORT that has the value of the Node port assigned:

```
> bash kubectl get service kubernetes-bootcamp -o jsonpath='{.spec.ports[0].nodePort}'
```

```
✓ run_shell_command kubectl get service kubernetes-bootcamp -o jsonpath='{.spec.ports[...
'30960'
```

```
PS C:\Users\LENOVO> echo "NODE_PORT=$NODE_PORT"
NODE_PORT=
```

Now we can test that the app is exposed outside of the cluster using curl, the IP address of the Node and the externally exposed port:

```
>> curl http://192.168.49.2:30960
wsl: 检测到 localhost 代理配置, 但未镜像到 WSL。NAT 模式下的 WSL 不支持 localhost 代理。
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.
```

Using labels

The Deployment created automatically a label for our Pod. With the describe

deployment subcommand you can see the name (the *key*) of that label:

```
> kubectl describe deployment
```

```
✓ run_shell_command kubectl describe deployment (描述所有部署的详细信息)
```

```
Name:                kubernetest-bootcamp
Namespace:            default
CreationTimestamp:    Fri, 10 Oct 2025 17:31:15 +0800
Labels:               app=kubernetest-bootcamp
Annotations:          deployment.kubernetest.io/revision: 1
Selector:             app=kubernetest-bootcamp
Replicas:             1 desired | 1 updated | 1 total | 1 available | 0
unavailable
StrategyType:        RollingUpdate
MinReadySeconds:      0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=kubernetest-bootcamp
  Containers:
    kubernetest-bootcamp:
      Image:   gcr.io/google-samples/kubernetest-bootcamp:v1
      Port:    <none>
      Host Port:  <none>
      Environment:  <none>
      Mounts:       <none>
      Volumes:      <none>
      Node-Selectors:  <none>
      Tolerations:    <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Available       True    MinimumReplicasAvailable
    Progressing     True    NewReplicaSetAvailable
  OldReplicaSets:  <none>
  NewReplicaSet:   kubernetest-bootcamp-658f6cbd58 (1/1 replicas created)
  Events:          <none>
```

use the kubectl get pods command with -l as a parameter, followed by the label values:

```
> kubectl get pods -l app=kubernetest-bootcamp
```

```
✓ run_shell_command kubectl get pods -l app=kubernetest-bootcamp (获取标签为 app=kubern...
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetest-bootcamp-658f6cbd58-f967h	1/1	Running	0	89m

Get the name of the Pod and store it in the POD_NAME environment variable:

```
> bash kubectl get pods -o jsonpath='{.items[*].metadata.name}'
```

```
✓ run_shell_command kubectl get pods -o jsonpath='{.items[*].metadata.name}' (获取所有...
```

```
'kubernetest-bootcamp-658f6cbd58-f967h'
```

```
> bash set POD_NAME=kubernetest-bootcamp-658f6cbd58-f967h
```

To apply a new label we use the label subcommand followed by the object type, object name

and the new label:

```
> bash kubectl label pods kubernetes-bootcamp-658f6cbd58-f967h version=v1
```

```
✓ run_shell_command kubectl label pods kubernetes-bootcamp-658f6cbd58-f967h version=v1...  
pod/kubernetes-bootcamp-658f6cbd58-f967h labeled
```

This will apply a new label to our Pod (we pinned the application version to the Pod), and we can check it with the describe pod command

```
> bash kubectl describe pods kubernetes-bootcamp-658f6cbd58-f967h
```

```
✓ run_shell_command kubectl describe pods kubernetes-bootcamp-658f6cbd58-f967h (
```

Name: kubernetes-bootcamp-658f6cbd58-f967h
Namespace: default
Priority: 0
Service Account: default
Node: minikube/192.168.49.2
Start Time: Fri, 10 Oct 2025 17:31:15 +0800
Labels: app=kubernetes-bootcamp
pod-template-hash=658f6cbd58
version=v1
Annotations: <none>
Status: Running
IP: 10.244.0.22
IPs: IP: 10.244.0.22
Controlled By: ReplicaSet/kubernetes-bootcamp-658f6cbd58
Containers:
 kubernetes-bootcamp:
 Container ID: docker://b4643640cb85c482d906987345c646f0305f5a789c95a7284dfbec03a98c601b
 Image: gcr.io/google-samples/kubernetes-bootcamp:v1
 Image ID: docker-pullable://gcr.io/google-samples/kubernetes-bootcamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f00e279c8fcc64af
 Port: <none>
 Host Port: <none>
 State: Running
 Started: Fri, 10 Oct 2025 17:31:35 +0800
 Ready: True
 Restart Count: 0
 Environment: <none>
 Mounts:
 /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-lxgdb (ro)
Conditions:
 Type Status
 PodReadyToStartContainers True
 Initialized True
 Ready True
 ContainersReady True
 PodScheduled True
Volumes:
 kube-api-access-lxgdb:
 Type: Projected (a volume that contains injected

```
data from multiple sources)
  TokenExpirationSeconds: 3607
  ConfigMapName: kube-root-ca.crt
  ConfigMapOptional: <nil>
  DownwardAPI: true
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute
op=Exists for 300s
              node.kubernetes.io/unreachable:NoExecute
op=Exists for 300s
Events: <none>
```

We see here that the label is attached now to our Pod. And we can query now the list of pods using the new label:

```
> kubectl get pods -l version=v1
```

✓ run_shell_command kubectl get pods -l version=v1 (获取标签为 version=v1 的 Pod)

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-658f6cbd58-f967h	1/1	Running	0	95m

Deleting a service

```
kubectl delete service -l app=kubernetes-bootcamp
```

✓ run_shell_command kubectl delete service -l app=kubernetes-bootcamp (删除标签为 app=... service "kubernetes-bootcamp" deleted

Scaling a Deployment

list Deployments

```
kubectl get deployments
```

✓ run_shell_command kubectl get deployments (获取所有部署的信息)

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
kubernetes-bootcamp	1/1	1	1	104m

see the ReplicaSet created by the Deployment

```
PS C:\Users\LENOVO> kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
kubernetes-bootcamp-658f6cbd58	1	1	1	107m

scale the Deployment to 4 replicas.

```
PS C:\Users\LENOVO> kubectl scale deployments/kubernetes-bootcamp --replicas=4
deployment.apps/kubernetes-bootcamp scaled
```

```
PS C:\Users\LENOVO> kubectl get deployments
NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
kubernetes-bootcamp                4/4      4              4            115m
```

check if the number of Pods changed

```
PS C:\Users\LENOVO> kubectl get pods -o wide
NAME                                READY    STATUS    RESTARTS    AGE    IP            NODE
kubernetes-bootcamp-658f6cbd58-5f65q 1/1      Running   0            3m2s   10.244.0.23   minikube
kubernetes-bootcamp-658f6cbd58-f967h 1/1      Running   0            118m   10.244.0.22   minikube
kubernetes-bootcamp-658f6cbd58-p4h6h 1/1      Running   0            3m2s   10.244.0.24   minikube
kubernetes-bootcamp-658f6cbd58-wwbh6 1/1      Running   0            3m2s   10.244.0.25   minikube
```

```
PS C:\Users\LENOVO> kubectl describe deployments/kubernetes-bootcamp
Name:                                kubernetes-bootcamp
Namespace:                           default
CreationTimestamp:                   Fri, 10 Oct 2025 17:31:15 +0800
Labels:                              app=kubernetes-bootcamp
Annotations:                         deployment.kubernetes.io/revision: 1
Selector:                           app=kubernetes-bootcamp
Replicas:                            4 desired | 4 updated | 4 total | 4 available | 0 unavailable
StrategyType:                        RollingUpdate
MinReadySeconds:                     0
RollingUpdateStrategy:               25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=kubernetes-bootcamp
  Containers:
    kubernetes-bootcamp:
      Image:          gcr.io/google-samples/kubernetes-bootcamp:v1
      Port:           <none>
      Host Port:      <none>
      Environment:    <none>
      Mounts:         <none>
      Volumes:        <none>
      Node-Selectors: <none>
      Tolerations:    <none>
Conditions:
  Type           Status  Reason
  ----           -
  Progressing    True    NewReplicaSetAvailable
  Available      True    MinimumReplicasAvailable
OldReplicaSets: <none>
NewReplicaSet:  kubernetes-bootcamp-658f6cbd58 (4/4 replicas created)
Events:
  Type           Reason             Age    From                      Message
  ----           -
  Normal         ScalingReplicaSet  3m26s  deployment-controller     Scaled up replica set kubernetes-bootcamp-658f6cbd58 from 1 to 4
```

check that the Service is load-balancing the traffic.

```
bash kubectl expose deployment kubernetes-bootcamp --type=LoadBalancer --port=8080

run_shell_command kubectl expose deployment kubernetes-bootcamp --type=LoadBalancer ...
service/kubernetes-bootcamp exposed
```

```
bash kubectl get service kubernetes-bootcamp -o
jsonpath='{.spec.ports[0].nodePort}'
```

```
✓ run_shell_command kubectl get service kubernetes-bootcamp -o jsonpath='{.spec.ports[...
'30759'
```

```
> bash set NODE_PORT=30759
```

```
✓ run_shell_command set NODE_PORT=30759 (将 NodePort 存储在 NODE_PORT 环境变量中)
```

do a curl to the exposed IP address and port.

```
PS C:\Users\LENOVO> curl 127.0.0.1:58664
```

```
StatusCode      : 200
StatusDescription : OK
Content         : Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-658f6cbd58-5f65q
                  | v=1
RawContent      : HTTP/1.1 200 OK
                  Connection: keep-alive
                  Transfer-Encoding: chunked
                  Content-Type: text/plain
                  Date: Fri, 10 Oct 2025 11:35:56 GMT
Forms           : {}
Headers         : {[Connection, keep-alive], [Transfer-Encoding, chunked], [Content-Type, text/
                  plain], [Date, Fri, 10 Oct 2025 11:35:56 GMT]}
Images          : {}
InputFields     : {}
Links           : {}
ParsedHtml      : mshtml.HTMLDocumentClass
RawContentLength : 84
                  Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-658f6...
```

Scale Down

```
PS C:\Users\LENOVO> kubectl scale deployments/kubernetes-bootcamp --replicas=2
deployment.apps/kubernetes-bootcamp scaled
```

The number of replicas decreased to 2.

```
PS C:\Users\LENOVO> kubectl get deployments
NAME                READY  UP-TO-DATE  AVAILABLE  AGE
kubernetes-bootcamp 2/2    2           2          125m
PS C:\Users\LENOVO> kubectl get pods -o wide
NAME                READY  STATUS      RESTARTS  AGE  IP           NODE
NOMINATED NODE     READINESS GATES
kubernetes-bootcamp-658f6cbd58-5f65q 1/1    Terminating  0      10m  10.244.0.23  minikube
kubernetes-bootcamp-658f6cbd58-f967h 1/1    Running       0      125m  10.244.0.22  minikube
kubernetes-bootcamp-658f6cbd58-p4h6h 1/1    Terminating  0      10m  10.244.0.24  minikube
kubernetes-bootcamp-658f6cbd58-wwbh6 1/1    Running       0      10m  10.244.0.25  minikube
```

Update the version of the app

```
C:\Users\LENOVO>kubectl get deployments
NAME                READY  UP-TO-DATE  AVAILABLE  AGE
kubernetes-bootcamp 2/2    2           2          14h

C:\Users\LENOVO>kubectl get pod
NAME                                READY  STATUS    RESTARTS  AGE
kubernetes-bootcamp-658f6cbd58-f967h 1/1    Running   1 (12m ago)  14h
kubernetes-bootcamp-658f6cbd58-wwbh6 1/1    Running   1 (12m ago)  12h

C:\Users\LENOVO>kubectl describe pods
Name:                kubernetes-bootcamp-658f6cbd58-f967h
Namespace:           default
Priority:              0
Service Account:     default
Node:                 minikube/192.168.49.2
Start Time:           Fri, 10 Oct 2025 17:31:15 +0800
Labels:               app=kubernetes-bootcamp
                     pod-template-hash=658f6cbd58
                     version=v1
Annotations:          <none>
Status:               Running
IP:                   10.244.0.29
IPs:                  IP: 10.244.0.29
Controlled By:        ReplicaSet/kubernetes-bootcamp-658f6cbd58
Containers:
  kubernetes-bootcamp:
    Container ID:      docker://d15164d49dca04b4f14303098eb8512736dba7e110e90fc5c5415f05d7d4a443
    Image:              gcr.io/google-samples/kubernetes-bootcamp:v1
```

update the image of the application to version 2

```
> kubectl set image deployments/kubernetes-bootcamp
kubernetes-bootcamp=docker.io/jocatalin/kubernetes-bootcamp:v2

✓ Shell kubectl set image deployments/kubernetes-bootcamp kubernetes-bootcamp=docker.io/jocatalin/kub
deployment.apps/kubernetes-bootcamp image updated
```

Check the status of the new Pods, and view the old one terminating with the get pods subcommand:

```
C:\Users\LENOVO>kubectl get pods
NAME                                READY  STATUS    RESTARTS  AGE
kubernetes-bootcamp-57cc954bb9-bf65p 1/1    Running   0          5m
kubernetes-bootcamp-57cc954bb9-c45z4 1/1    Running   0         4m55s
```

Verify an update

Create an environment variable called NODE_PORT that has the value of the Node port assigned:

```
✓ Shell kubectl get service kubernetes-bootcamp -o jsonpath="{.spec.ports[0].nodePort}"
"30759"

✓ Shell set NODE_PORT=30759 && cmd /c echo NODE_PORT=%NODE_PORT%
```

```
Kubernetes Bootcamp App Started At: 2025-10-11T00:16:44.279Z | Running On:
kubernetes-bootcamp-57cc954bb9-bf65p
```

You can also confirm the update by running the rollout status subcommand:

```
C:\Users\LENOVO>kubectl rollout status deployments/kubernetes-bootcamp
deployment "kubernetes-bootcamp" successfully rolled out
```

To view the current image version of the app, run the describe pods subcommand:

```
Image:          docker.io/jocatalin/kubernetes-bootcamp:v2
Image ID:
docker-pullable://jocatalin/kubernetes-bootcamp@sha256:fb1a3ced00cecf
c1f83f18ab5cd14199e30adc1b49aa4244f5d65ad3f5feb2a5
```

Roll back an update

deploy an image tagged with v10:

```
Container ID:
Image:          gcr.io/google-samples/kubernetes-bootcamp:v10
Image ID:
```

To roll back the deployment to your last working version, use the rollout undo subcommand:

```
C:\Users\LENOVO>kubectl rollout undo deployments/kubernetes-bootcamp
deployment.apps/kubernetes-bootcamp rolled back
```

Use the get pods subcommand to list the Pods again:

```
C:\Users\LENOVO>kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-57cc954bb9-bf65p	1/1	Running	0	58m
kubernetes-bootcamp-57cc954bb9-c45z4	1/1	Running	0	58m
kubernetes-bootcamp-677ff875c4-q747l	0/1	ErrImagePull	0	40s

check the image deployed on the running Pods

```
2c8
Image:          docker.io/jocatalin/kubernetes-bootcamp:v2
Image ID:
docker-pullable://jocatalin/kubernetes-bootcamp@sha256:fb1a3ced00cecf
1f83f18ab5cd14199e30adc1b49aa4244f5d65ad3f5feb2a5
Port:          <none>
```